




CivicMind Jing-Zhang | Governance as Infrastructure

Rules, public review and trusted services form one accountable loop



Xiaoyue River testing wing

Low-risk, reversible community tests

Zhongguancun service wing

Turn validated methods into standards, tools and talent

02 / One body, two wings, three layers · Governance loop

The Jing-Zhang heritage governance spine coordinates three functional layers; two wings connect real-world testing and professional service export.



Three Key Areas | Three Civic Experiences

Each area pairs a public action, spatial toolkit and safeguard

A1 Governance Loop

SEE HOW RULES ARE MADE

Knowledge gallery • simulation room
Compliance test court
Safeguard: models never replace statutory duty



A2 Sensing Eye

CHALLENGE AND REACH A PERSON


Public review hall • accessible navigation
Co-design court
Safeguard: no face or route collection



A3 Trust Bridge


UNDERSTAND PERMITTED DATA USE

Trusted exchange • scenario hall
Legal-service referral
Safeguard: key decisions stay human-led



04/ **Six users · Six real service journeys**

Personas test whether urban-agent services remain equal, legible and transferable to a responsible human under different needs and risks.



05 / 11 AI + governance scenarios

A shared service chain: Discover → Assess → Test → Review → Update. Every scenario defines data boundaries, Human Review and exit conditions.

Civic-agent governance Enterprise ecosystem AI + public services

S01 · DISCOVER
City Sensing Grid
Aggregates public complaints, hotlines, open data, sensors and grid reports to identify urban issues.
Residents · regulators | A2 | High-risk alerts require review
Civic-agent governance

S02 · ASSESS
Policy Simulation Sandbox
Compares policy impacts, uncertainty and distributional risk without issuing administrative conclusions.
Policy teams · researchers | A1 | Accountable sign-off
Civic-agent governance

S03 · COLLABORATE
Multi-agent Joint Review
Matches evidence, responsibilities and workflows across departments; staff approve the joint plan.
Regulators | A1 | Human approval before action
Civic-agent governance

S04 · LEARN
Governance Knowledge Base
Structures rules, cases, workflows and reviewed outcomes for traceable assisted analysis.
Frontline staff · researchers | A1 | Knowledge review
Civic-agent governance

S05 · SUPERVISE
Public Review Desk
Explains grounds, accountable owner and correction route; accepts challenge and fairness complaints.
Residents · underserved groups | A2 | Appeal and correction
Civic-agent governance

S06 · OPEN
AI Scenario Marketplace
Publishes real needs, admission rules, protocols, acceptance criteria and exit conditions.
Enterprises · government | A3 / Xiaoyue | Scenario review
Enterprise ecosystem

S07 · TRUST
Trusted Data Exchange Hub
Uses catalogues, purpose-bound access, privacy computing and audit logs for accountable use.
Enterprises · departments · universities | A3 | Permission approval
Enterprise ecosystem

S08 · COMPLY
AI Compliance Accelerator
Offers regulatory search, risk self-check, expert guidance, controlled tests and monitoring.
AI firms · founders | A1 | Expert review; no automatic certification
Enterprise ecosystem

S09 · SERVE
AI Civic Navigator
Recognises service intent, provides accessible multimodal guidance and transfers complex cases to staff.
Residents · visitors | A2 | Human transfer
AI + public services

S10 · HEALTH
Community Health Sentinel
Detects aggregate service pressure to support community resource readiness; no individual diagnosis.
Residents · health workers | Xiaoyue | Professional review
AI + public services

S11 · REMEDY
AI Legal Aid Station
Supports legal-information search and document preparation, then refers cases to qualified legal aid.
Underserved groups · SMEs | A3 | Lawyer review
AI + public services

06 / Urban-agent supervision & Human Review dashboard

A single conceptual entry point for operating status, risk, public challenge and accountable human judgement. It does not issue administrative conclusions.

Urban-Agent Governance Supervision Desk

Concept interface · Aggregated data · Tiered permission

• System operating normally

11-scenario operating health



Pending Human Review




Risk & permission

- ⚠ 1 authorisation near expiry
- ⚠ 2 false-positive reviews due
- ✓ 0 unauthorised calls

01 Read-only monitoring



02 Human assessment



03 Accountable sign-off




04 Audit & review

Illustrative values show information architecture only, they are not live operating data or evidence of government deployment.

07 / Blue-green public space & governance-service network

Walking routes, ecological corridors, heritage spaces and service nodes create a civic interface that is reachable, understandable and able to provide human help.

A reachable, stoppable and assisted public-service chain



From place to service

Spine: Jing-Zhang culture, slow mobility and governance literacy.

Nodes: Public review, accessible guidance and staffed service.

Continuous walking

Walking, cycling and transit information connect the heritage spine.

Accessible service chain

Voice, large text, tactile routes, low counters and visible human help.

Minimum sensing

Collect only the least data necessary for an authorised public purpose; high-risk alerts require review and logging.

From Concept to Delivery | Three Evidence Gates

Public value and spatial logic lead; technical evidence supports the decision

01 SPATIAL BASIS

Label provisional limits
Keep topology complete
Recalculate with official data

02 ACCOUNTABILITY

Name users and owners
Set data limits and review
Define exit conditions

03 DELIVERY GATE

Start with reversible pilots
Scale after value and safety tests
Stop or exit if tests fail

VERIFIABLE NOW

3 key areas · 11 scenarios · 6 personas · 3 landmarks

REQUIRES PROFESSIONAL INPUT

FAR, height, density, statutory green ratio, setbacks, ownership and utilities

From one service failure to a correction receipt

Synthetic desktop situation: an accessible passage suggested by A2 navigation is unavailable. Provide staffed help first, pause disputed information, then verify the space, correct the content and obtain independent review.

A2 · Discover and take over

S09 navigation / S05 public review. Staff assist first, record the issue and old version, and require no disability proof, face data, diagnoses or movement traces.

A1 · Verify and correct

S03 collaboration / S04 knowledge. Site and content reviewers verify conditions and versions separately. Keep the pause if evidence is missing; refer product defects to A3.

A2 · Receipt and further review

Restore only after field verification, staffed access, retest and sign-off. Supply a paper or accessible receipt with reasons and a further-review route.

Concept and synthetic record only. No live connection or field measurements. Existing dashboard numbers such as 07 remain illustrative and are not counts derived from this sample.

Regional collaboration through accountable handovers

Consent, authority and verifiable outputs precede collaboration. The Beiwei Community interface adds an early-product test entry; other interfaces support scientific testing, public explanation, remediation and cross-region adaptation.

Beiwei Community / A1/A2/A3

Willing teams supply public product issues and synthetic samples. A1 co-tests rules, A2 tests explanations and staffed referrals, and A3 retests corrections, returning versioned gaps and exit receipts.

Scientific methods / local review

The Future Science City interface supports controlled rule simulation; the Huairou interface explores scientific tools and public explanation. Reviewed outputs return to A1 or A2 respectively.

Industry and regions / reusable methods

The E-Town interface co-tests enterprise services; Jing-Jin-Ji interfaces examine local applicability. The Zhongguancun wing shares methods; the Xiaoyue River wing receives controlled feedback.

These are proposed interfaces, not commitments from districts, institutions or teams. Do not start without accountability, permission, maintenance and resources.

Three pilots: from workflow to acceptance evidence

Freeze inputs and rule versions, retaining failed and insufficient-evidence cases. Rule checks do not assess AI quality or establish that a live pilot has passed.

S03 Multi-agent collaboration

Input: synthetic issues and human responsibility baseline. Inject conflict, missing evidence and absent sign-off. Route conflicts to staff and prevent progression without evidence and sign-off.

S06 Scenario marketplace

Input: synthetic applications and admission checklist. Inject expired permission, withdrawal and missing maintenance ownership. Reject or exit; require new review, with no certification or procurement.

S10 Community health sentinel

Input: synthetic aggregate demand and a professional alert baseline. Inject small samples, noise and unavailable staff. Insufficient evidence cannot pass; pause without staff; no individual diagnosis.

Rule source and results: `visual/assets/pilot-rule-source.json` and `pilot-rule-results.json`. Extract the Python source as instructed in the JSON. Checks cover Boolean gates and stop rules, not real performance or medical outcomes.