




AI-generated conceptual visualization. It does not represent an official boundary, an approved construction plan, or an implementation commitment.

CivicMind Jing-Zhang | Governance as Infrastructure

Rules, public review and trusted services form one accountable loop



Xiaoyue River testing wing

Low-risk, reversible community tests

Zhongguancun service wing

Turn validated methods into standards, tools and talent

02 / One body, two wings, three layers · Governance loop

The Jing-Zhang heritage governance spine coordinates three functional layers; two wings connect real-world testing and professional service export.



Three Key Areas | Three Civic Experiences

Each area pairs a public action, spatial toolkit and safeguard

A1 Governance Loop

SEE HOW RULES ARE MADE

Knowledge gallery • simulation room

Compliance test court

Safeguard: models never replace statutory duty



A2 Sensing Eye

CHALLENGE AND REACH A PERSON

Public review hall • accessible navigation

Co-design court

Safeguard: no face or route collection



A3 Trust Bridge

UNDERSTAND PERMITTED DATA USE

Trusted exchange • scenario hall

Legal-service referral

Safeguard: key decisions stay human-led



04/ Six users · Six real service journeys

Personas test whether urban-agent services remain equal, legible and transferable to a responsible human under different needs and risks.

R Ms Wang | Local resident

A neighbourhood health alert appears. She receives a plain-language explanation, chooses human advice, and may inspect the evidence and response status.

Health Sentinel → Civic Navigation → Public Review

F Li | Young founder

He seeks a real-world pilot, reads admission and exit rules, enters controlled testing, and receives a remediation note if the test fails.

Scenario Access → Compliance Support → Trusted Data

E Zhang | Enterprise data lead

She requests a public-data catalogue and can verify purpose, permission, term and audit duties, with a clear route for disputes or changes.

Data Catalogue → Permission Review → Audit Trail

A Prof Chen | Researcher

She tests policy options with de-identified data in a controlled environment; accountable departments and ethics reviewers confirm the research boundary.

Problem Brief → Policy Simulation → Result Feedback

V Mr Zhao | Short-stay visitor

He uses continuous wayfinding to understand Jing-Zhang culture, reach transit and open experience points, and transfer quickly to staffed service.

Gateway Guide → Culture Spine → Human Referral

I Zhou | Visually impaired resident

He uses voice, tactile routes and low counters; where an algorithm may cause unfair impact, he can seek legal aid and Human Review.

Accessible Guide → Legal Aid → Fairness Review

05 / 11 AI + governance scenarios

A shared service chain: Discover → Assess → Test → Review → Update. Every scenario defines data boundaries, Human Review and exit conditions.

Civic-agent governance

Enterprise ecosystem

AI + public services

S01 · DISCOVER

City Sensing Grid

Aggregates public complaints, hotlines, open data, sensors and grid reports to identify urban issues.

Residents · regulators | A2 | High-risk alerts require review

Civic-agent governance

S02 · ASSESS

Policy Simulation Sandbox

Compares policy impacts, uncertainty and distributional risk without issuing administrative conclusions.

Policy teams · researchers | A1 | Accountable sign-off

Civic-agent governance

S03 · COLLABORATE

Multi-agent Joint Review

Matches evidence, responsibilities and workflows across departments; staff approve the joint plan.

Regulators | A1 | Human approval before action

Civic-agent governance

S04 · LEARN

Governance Knowledge Base

Structures rules, cases, workflows and reviewed outcomes for traceable assisted analysis.

Frontline staff · researchers | A1 | Knowledge review

Civic-agent governance

S05 · SUPERVISE

Public Review Desk

Explains grounds, accountable owner and correction route; accepts challenge and fairness complaints.

Residents · underserved groups | A2 | Appeal and correction

Civic-agent governance

S06 · OPEN

AI Scenario Marketplace

Publishes real needs, admission rules, protocols, acceptance criteria and exit conditions.

Enterprises · government | A3 / Xiaoyue | Scenario review

Enterprise ecosystem

S07 · TRUST

Trusted Data Exchange Hub

Uses catalogues, purpose-bound access, privacy computing and audit logs for accountable use.

Enterprises · departments · universities | A3 | Permission approval

Enterprise ecosystem

S08 · COMPLY

AI Compliance Accelerator

Offers regulatory search, risk self-check, expert guidance, controlled tests and monitoring.

AI firms · founders | A1 | Expert review; no automatic certification

Enterprise ecosystem

S09 · SERVE

AI Civic Navigator

Recognises service intent, provides accessible multimodal guidance and transfers complex cases to staff.

Residents · visitors | A2 | Human transfer

AI + public services

S10 · HEALTH

Community Health Sentinel

Detects aggregate service pressure to support community resource readiness; no individual diagnosis.

Residents · health workers | Xiaoyue | Professional review

AI + public services

S11 · REMEDY

AI Legal Aid Station

Supports legal-information search and document preparation, then refers cases to qualified legal aid.

Underserved groups · SMEs | A3 | Lawyer review

AI + public services

06 / **Urban-agent supervision & Human Review dashboard**


A single conceptual entry point for operating status, risk, public challenge and accountable human judgement. It does not issue administrative conclusions.



07 / Blue-green public space & governance-service network

Walking routes, ecological corridors, heritage spaces and service nodes create a civic interface that is reachable, understandable and able to provide human help.

A reachable, stoppable and assisted public-service chain



From place to service

Spine: Jing-Zhang culture, slow mobility and governance literacy.

Nodes: Public review, accessible guidance and staffed service.

Continuous walking

Walking, cycling and transit information connect the heritage spine.

Accessible service chain

Voice, large text, tactile routes, low counters and visible human help.

Minimum sensing

Collect only the least data necessary for an authorised public purpose; high-risk alerts require review and logging.

From Concept to Delivery | Three Evidence Gates

Public value and spatial logic lead; technical evidence supports the decision

01 SPATIAL BASIS

Label provisional limits
Keep topology complete
Recalculate with official data

02 ACCOUNTABILITY

Name users and owners
Set data limits and review
Define exit conditions

03 DELIVERY GATE

Start with reversible pilots
Scale after value and safety tests
Stop or exit if tests fail

VERIFIABLE NOW

3 key areas · 11 scenarios · 6 personas · 3 landmarks

REQUIRES PROFESSIONAL INPUT

FAR, height, density, statutory green ratio, setbacks, ownership and utilities

From one service failure to a correction receipt

Synthetic desktop situation: an accessible passage suggested by A2 navigation is unavailable. Provide staffed help first, pause disputed information, then verify the space, correct the content and obtain independent review.

A2 · Discover and take over

S09 navigation / S05 public review. Staff assist first, record the issue and old version, and require no disability proof, face data, diagnoses or movement traces.

A1 · Verify and correct

S03 collaboration / S04 knowledge. Site and content reviewers verify conditions and versions separately. Keep the pause if evidence is missing; refer product defects to A3.

A2 · Receipt and further review

Restore only after field verification, staffed access, retest and sign-off. Supply a paper or accessible receipt with reasons and a further-review route.

Concept and synthetic record only. No live connection or field measurements. Existing dashboard numbers such as 07 remain illustrative and are not counts derived from this sample.

Regional collaboration through accountable handovers

Consent, authority and verifiable outputs precede collaboration. The Beiwei Community interface adds an early-product test entry; other interfaces support scientific testing, public explanation, remediation and cross-region adaptation.

Beiwei Community / A1/A2/A3

Willing teams supply public product issues and synthetic samples. A1 co-tests rules, A2 tests explanations and staffed referrals, and A3 retests corrections, returning versioned gaps and exit receipts.

Scientific methods / local review

The Future Science City interface supports controlled rule simulation; the Huairou interface explores scientific tools and public explanation. Reviewed outputs return to A1 or A2 respectively.

Industry and regions / reusable methods

The E-Town interface co-tests enterprise services; Jing-Jin-Ji interfaces examine local applicability. The Zhongguancun wing shares methods; the Xiaoyue River wing receives controlled feedback.

These are proposed interfaces, not commitments from districts, institutions or teams. Do not start without accountability, permission, maintenance and resources.

Three pilots: from workflow to acceptance evidence

Freeze inputs and rule versions, retaining failed and insufficient-evidence cases. Rule checks do not assess AI quality or establish that a live pilot has passed.

S03 Multi-agent collaboration

Input: synthetic issues and human responsibility baseline. Inject conflict, missing evidence and absent sign-off. Route conflicts to staff and prevent progression without evidence and sign-off.

S06 Scenario marketplace

Input: synthetic applications and admission checklist. Inject expired permission, withdrawal and missing maintenance ownership. Reject or exit; require new review, with no certification or procurement.

S10 Community health sentinel

Input: synthetic aggregate demand and a professional alert baseline. Inject small samples, noise and unavailable staff. Insufficient evidence cannot pass; pause without staff; no individual diagnosis.

Rule source and results: `visual/assets/pilot-rule-source.json` and `pilot-rule-results.json`. Extract the Python source as instructed in the JSON. Checks cover Boolean gates and stop rules, not real performance or medical outcomes.